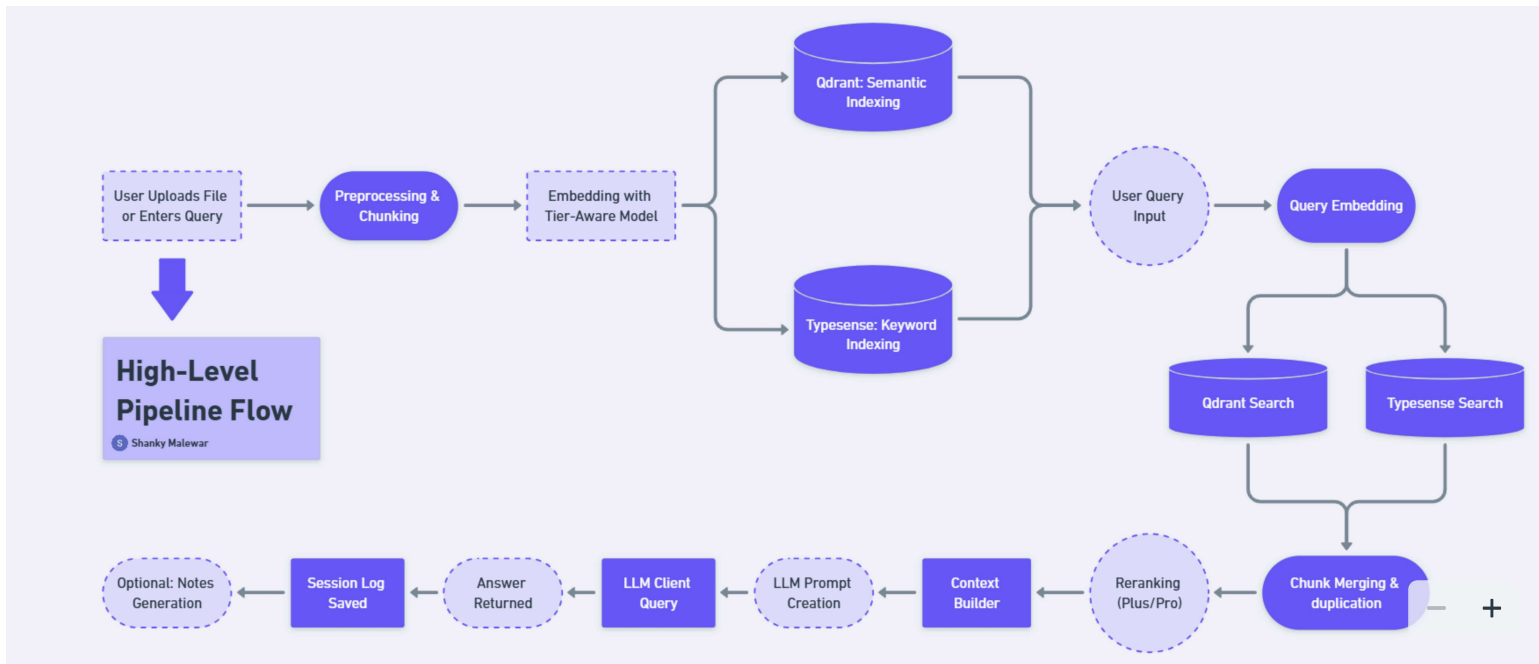




# Asklyne RAG Architecture — Detailed System Design!

Asklyne’s core AI capabilities are powered by a custom `Retrieval-Augmented Generation (RAG)` system that goes far beyond traditional static LLM chat. It allows Asklyne to process and reason over user-uploaded `PDFs, images, code, and text documents` — providing meaningful, context-aware answers, summaries, and learning notes.

This document outlines how Asklyne’s end-to-end RAG pipeline functions across multiple user tiers (`free`, `plus`, `pro`) and input types (`text`, `code`, `notes`).



## ➡ 1 Upload & Preprocessing (/upload-file)

When a user uploads a file, the following steps take place:

Mode	Processing Strategy
Notes(images)	OCR using <code>Tesseract</code> to extract visible text
Code	Parsed using custom <code>extract_code_from_*</code> logic
Text(.pdf .txt)	Extracted using <code>pdfplumber</code> (for PDFs) or UTF-8 decode

## ➡ 2 Chunking (Token-Aware Splitting)

To ensure the LLM can work within token limits, raw content is split into semantically meaningful chunks using the `Chunker` class:



### Chunking Logic

- **Text Mode:** Sentence-based splitting with overlap (`max_tokens = 480`, `overlap = 80`)

- **Code Mode:** Function/class-based splitting using regex anchors on def and class
- **Overlap** ensures contextual continuity between adjacent chunks for more coherent retrieval

### ➡ 3 Embedding the Chunks

Chunks are transformed into high-dimensional vectors using tier-specific sentence encoders from sentence-transformers. This is handled by the Embedder class.

Tier	Text Embedding	Code Embedding
Free	BAAI/bge-large-en-v1.5	microsoft/codebert-base
Plus	BAAI/bge-large-en-v1.5	Salesforce/codet5-base
Pro	intfloat/multilingual-e5-large	microsoft/graphcodebert-base

### ➡ 4 Storage in Dual-Engine Vector DBs

Asklyne indexes chunks in **two complementary retrieval systems**:

#### 📦 Qdrant (Semantic Search)

- Stores vectors for **cosine similarity search**
- Collections per (tier, mode): e.g. asklyne\_chunks\_plus\_code
- Embeddings are stored using PointStruct along with metadata (session\_id, mode, tier, text)
- Supports filtering by session for scoped queries

#### 🔍 Typesense (Keyword Search)

- Full-text index of each chunk for **keyword-based matching**
- Schema includes text, tier, mode, session\_id
- Complements Qdrant by catching lexical matches not captured by vector similarity

This **hybrid retrieval strategy** ensures that both conceptual and literal matches are surfaced.

### ➡ 5 Query Handling

When a user submits a query, Asklyne executes an orchestrated multi-step retrieval & generation process:

#### 🔄 a. Retriever

- Embeds the query using the same encoder as the stored chunks
- Fetches top `k` relevant chunks from:
  - **Qdrant** (semantic match)
  - **Typesense** (keyword match)
- Deduplicates results by chunk text

#### 📊 b. Reranker (Plus/Pro Only)

- For higher tiers, chunks are reranked using `cross-encoder/ms-marco-MiniLM-L-6-v2`
- Each chunk receives a `score` and is sorted accordingly



### c. Context Builder

- Merges reranked top chunks into a prompt-safe block
- Uses 90% of tier's token limit to preserve headroom
- Format: `Chunk1\n---\nChunk2\n---\n...`



6

## Notes Generation ( `/generate-notes` + `/generate-notes-pdf` )

Asklyne enables users to turn entire chat sessions into **clean, structured notes**:



### Features

- Generates notes using LLM prompting
- Markdown → HTML → PDF conversion
- Templated prompt includes:
  - Headings
  - Bullets
  - Concept highlights



### Modes

- `full` : Includes both Q & A
- `response_only` : Includes only answers
- `custom` : User-defined focus (e.g., "only summarize Python code explanations")



### Tier-Aware Intelligence Routing

Every layer — chunker, embedder, reranker, LLM, note model — is **dynamically routed based on user tier**, enabling:

- Lower costs for casual learners
- Powerful tools for researchers & devs
- Custom logic per mode ( `text` , `code` , `notes` )

This **modular and scalable architecture** allows Asklyne to balance cost and quality across diverse users.